



Fintech Daily · AI news agent

Every morning, 7 fact-checked fintech stories for professionals in Chinese, English and Russian. Automated from 54 sources to an on-the-minute post, with a human accountable for what goes out.

AI PRODUCT CASE · SOLO PROJECT

Product owner · built end to end

Sep 2026 – present · live

08:30 Singapore · @fintech_daily_zh 08:30 Moscow · @fintech_daily_ru 07:30 UTC · @fintech_daily
7 stories a day · week in review on Mondays

45 /45

issues published on time, none missed

0 min

publishing delay, down from 11.8 min; 33/33 on the minute

15%

of AI-written entries fixed by guardrails before publishing (46/306)

74%

of problems caught by deterministic code, not by an LLM

\$0

monthly running cost, all within free tiers

The problem PROBLEM

Fintech news is scattered across dozens of outlets, newsletters and regulators, mostly in English. A human-edited daily in three languages is expensive. Letting an LLM write finance news is risky: one wrong amount, or "pending" written as "approved", and readers stop trusting the channel.

How do you get an unstable model that hallucinates to publish trustworthy content on time, every day, on a zero budget?

Trust first

every fact traceable to the source

No repeats

yesterday's event never returns

On time is a promise

says 08:30, posts at 08:30

Zero cost

sustainable on free tiers

A real digest entry OUTPUT

@fintech_daily · 2026-10-08

story 1 of 7

Stripe and FedEx partner on SMB lending

Stripe and FedEx are teaming up to offer capital to smaller businesses by merging operational and financial data.


→ **Why it matters:** Combining logistics and financial metrics can unlock alternative underwriting signals for SMB credit.

Payments Dive

#lending #B2B

Multi-agent pipeline ARCHITECTURE

code LLM human



Key product decisions DECISIONS

- Route models by data, not by instinct**
An audit showed every new repeat came on days a lightweight fallback model did the selection → **wait for the main model only at selection**; other steps fall back and guardrails hold quality; a second repeat check added. 5 repeats caught in 3 days.
- On time is a product promise**
The channel said 08:30, posts went out around 08:43 → an external scheduler wakes the bot on the minute; every run is idempotent. Budgeted the free tier first and removed ~330 wasted minutes a month. **Delay 11.8 → 0 min.**
- Autonomy for the AI, a veto for the human**
Auto (veto) and manual (approve) modes with a 20-minute veto window. "Rewrite" may only use the original sources and re-runs every guardrail; **an edit that fails keeps the old text.**
- Launch is not the finish line**
A week after a fix that passed its tests, every post was audited again and a new root cause surfaced. **In LLM systems, passing tests is not the same as working in production.**

Four guardrail layers GUARDRAILS

Grounding	grounding
written from full text, with verbatim quotes	
Deterministic checks	34 / 46
numbers and sources must match	
Cross-model review	12 / 46
LLM-as-judge: status, roles, terms	
Human veto	human
one tap stops a post	
Evals & metrics: offline checks (0 model calls) · blind A/B LLM-as-judge on its own quota · every issue logs guardrail hits, fallbacks and token cost.	
Deliberately not built: long-form explainers. Short entries already need fixes 15% of the time; long texts multiply hallucination risk and review cost.	

AI-native way of working

Built with Claude Code as the AI coding tool. **I owned** problem discovery, requirements and acceptance criteria, trade-off decisions, PR review and post-launch audits; **the AI tool** proposed designs, wrote the code and the tests.

Engineering & stack

2 days to MVP · 113 unit tests + CI
Python 3.12 · Gemini Flash routed per step
GitHub Actions · Telegram Bot API

